

Improving Softmax Regression on MNIST

Last update: November 9, 2024



fit@hcmus

KHOA CÔNG NGHỆ THÔNG TIN
TRƯỜNG ĐẠI HỌC KHOA HỌC TỰ NHIÊN

Stochastic Gradient Descent (SGD)

Error function on training data

$$E_{train}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N e(h_{\mathbf{w}}(\mathbf{x}^{(n)}), y^{(n)})$$

Minimize $E_{train}(\mathbf{w})$

Gradient Descent (GD):

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \nabla E_{train}(\mathbf{w})$$

$$\nabla E_{train}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \nabla e(h_{\mathbf{w}}(\mathbf{x}^{(n)}), y^{(n)})$$

To take a step, GD needs to go through N training examples

→ Slow when N is large

Minimize $E_{train}(\mathbf{w})$

Stochastic Gradient Descent (SGD): let's use a subset of B ($B \ll N$) examples — called a *mini-batch* — to estimate $\nabla E_{train}(\mathbf{w})$

Q: How should we choose B examples from N examples?

A: Choose randomly

Q: What is the effect of B on the quality of estimating $\nabla E_{train}(\mathbf{w})$?

A: The larger B , the better quality of estimating $\nabla E_{train}(\mathbf{w})$

Minimize $E_{train}(\mathbf{w})$

SGD (left) vs GD (right)




Image source: <https://datascience.stackexchange.com/a/94355>

Minimize $E_{train}(\mathbf{w})$

Q: So, with a step of SGD, we do it **faster** than GD, but we pay the price of **worse quality**. Is it worth it?

A: Yes. In reality, people often see that: with the same starting point and the same amount of time, SGD runs much more steps than GD and **ends up with much smaller E_{train} than GD**

Minimize $E_{train}(\mathbf{w})$


Sketch of SGD implementation:

1. Initialize \mathbf{w}
2. Repeat until termination criteria are satisfied:
 - a. Shuffle the order of training examples
 - b. For mini-batch $b = 1, \dots, \frac{N}{B}$:

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{1}{B} \sum_{n=(b-1)B+1}^{bB} \nabla e(h_{\mathbf{w}}(\mathbf{x}^{(n)}), y^{(n)})$$


Deslanting digit image

Origin



I (input image)

Origin




O (output image)

Q: $O[r, c] = I[r, ?]$

Hint: $?$ can be computed from c , $\text{center_r} - r$, $\tan \alpha$

A: $?$ = $c + (\text{center_r} - r) \times \tan \alpha$


How to find the green point (center_r) and the violet line ($\tan \alpha$)?



A data point (a white point) represents row index and column index of a pixel corresponding to pen stroke

The green point is simply the average of all data points!

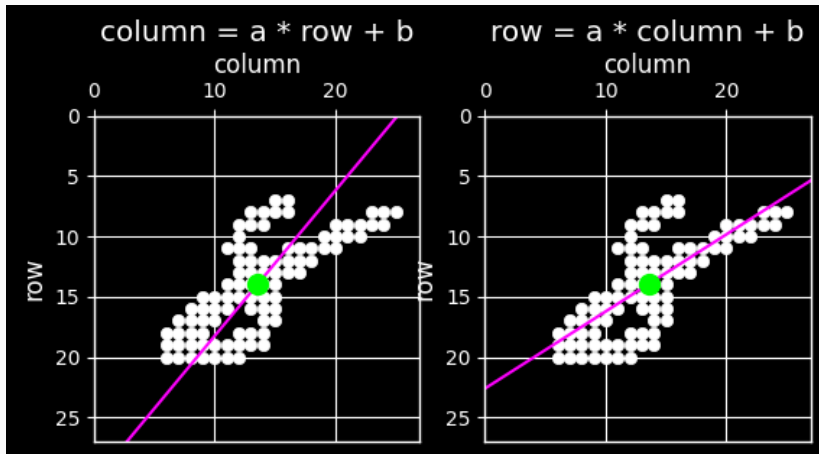
How to find the green point (center_r) and the violet line ($\tan \alpha$)?



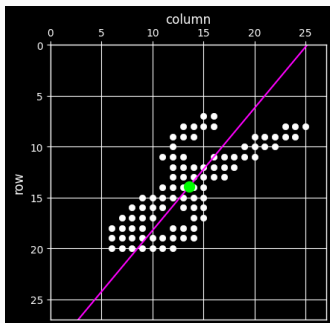
A data point (a white point) represents row index and column index of a pixel corresponding to pen stroke

The violet line is the best-fit line for all data points!

Q: Which is the form of this line: (1) row = $a \times \text{column} + b$,
or (2) column = $a \times \text{row} + b$?



How to find the green point (center_r) and the violet line ($\tan \alpha$)?




A data point (a white point) represents row index and column index of a pixel corresponding to pen stroke

The violet line is the best-fit line for all data points!

Q: Which is the form of this line: (1) row = $a \times \text{column} + b$,
or (2) column = $a \times \text{row} + b$?

A: (2)

How to find the green point (center_r) and the violet line ($\tan \alpha$)?



A data point (a white point) represents row index and column index of a pixel corresponding to pen stroke

The violet line is the best-fit line for all data points!

Q: Assume we have found the best-fit line “column = $a \times$ row + b ”. What is the formula to compute $\tan \alpha$ from a ?

A: $\tan \alpha = -a$